

Data Organization and Processing

File Organizations

David Hoksza, Petr Škoda
<http://siret.cz/hoksza>

Overview

- Terminology
 - records, files, ...
- Querying files
- File organizations
 - types of file organizations
 - implementation of file operations

Data organization - terminology

- **Database** (*databáze*)
 - collection of related data (named files) in secondary memory
 - elements are related
- **File** (*soubor*)
 - named collection of records
- **Record** (*záznam*)
 - representation of an application object (person, car, flat, ...)
 - **set of fields** (*pole*)
 - elementary unit of data
 - content provided by a user or a program
 - fixed or variable length
 - person – name, occupation, height, sex, ...
 - car – license plate, type, color, ...
 - **field F_i** together with its **domain $\text{dom}(F_i)$** , i.e. pair $F_i : \text{dom}(F_i)$, is called **attribute** (*atribut*)
 - domains: INT, FLOAT, STRING, ...

Data record

- **Logical**
 - attribute set
- **Physical**
 - **physical representation** of a logical record of size R (bytes) **on the medium**
 - contains **additional metadata** (such as record delimiters, etc.)
 - **Records** are stored **in blocks** of size B (bytes)
 - transferred between primary and secondary memory

Data record types

- **Fixed** length
 - **file header** contains **number of records** and **length of each field**
 - record can be accessed using the record number
- **Variable** length
 - variable length of the attributes (e.g. name, description, ...)
 - when **similar objects relate to a nonuniform set of data** (e.g. for different types of employees different attributes are relevant)
 - **optional** attributes (e.g. product picture)
 - **attributes holding records** (e.g. an order with multiple items or employee with multiple phone numbers)

Record blocking (1)

- **Blocking factor** (*blokovací faktor*)
 b

- number of records in a block

- $\left\lfloor \frac{B}{R} \right\rfloor$ **ratio**

- B ... block size
- R ... record size

- Basic division based on blocking

- **non-blocked** records
 - 1 record fits 1 block
- **blocked** records
 - N records fit 1 block
- **overflowed** records
 - 1 record fits N blocks

Record blocking (2)

Fixed blocking

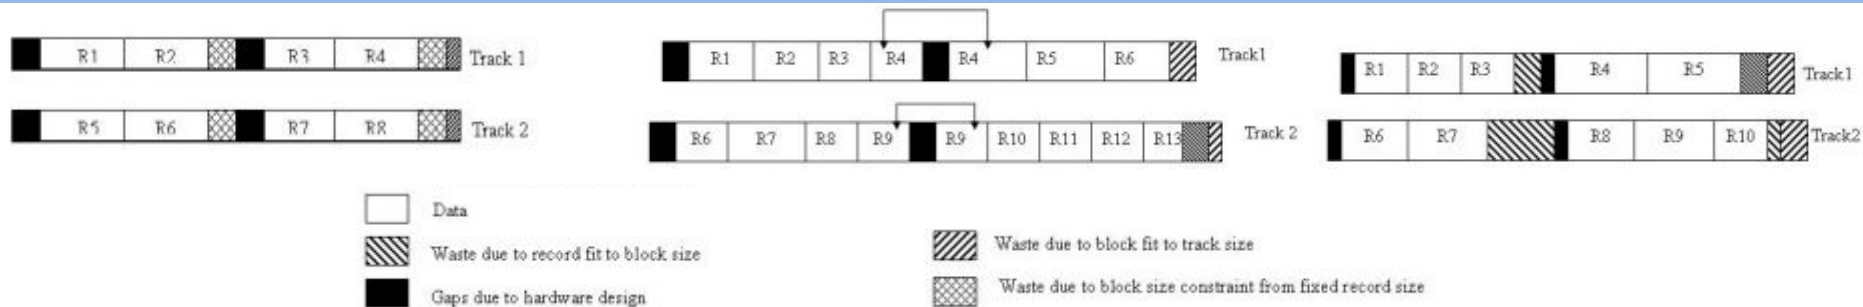
- fixed-length records
- possible internal fragmentation
 - unused space

Variable-length spanned blocking

- variable-length records
- no unused space
- continuation indicated by a pointer to the next block
- hard to implement
- need more time to read records in 2 blocks

Variable-length unspanned blocking

- variable-length records
- no spanning
- can show high internal fragmentation



Files

- Storage in secondary memory
 - **reading** from a file requires its contents to be read into main memory **in blocks**
 - when **modifying** a file its content needs to be **read** into main memory, **modified** and changes are then **written back** to the secondary memory
- **Homogenous**
 - stores fixed-length records of the same type
 - declared as $S(A_1: D_1, \dots A_N: D_N)$
- **Non-homogenous**
 - stores either variable-length records or records of different types

Operations over files

- **Formation/Termination**

- **CREATE, REMOVE**

- builds or removes auxiliary structures

- **Modification**

- **INSERT**

- inserts a new record into a file

- **UPDATE**

- updates an existing record

- **DELETE**

- removes (or invalidates) an existing record

- **Querying**

- **FIND**

- finds a record in a file corresponding to the query key

- **FETCH**

- gets a record from a file (secondary memory) into main memory based on specified query conditions (query key)

- **Maintenance**

- **REORGANIZE/REBUILD**

- due to the efficiency reasons not all changes are directly projected into the underlying file organization

Querying files

- **One-dimensional queries**

- querying according to only one attribute
 - fetch employees with age > 35
 - fetch cars with color = 'red'

- **Multi-dimensional queries**

- querying according to multiple attribute
- **total match** (*na úplnou shodu*)
 - the values of all attributes entered
- **partial match** (*na částečnou shodu*)
 - only values of some attributes entered
- **total interval match** (*na úplnou intervalovou shodu*)
 - for each attribute the interval of values entered
- **partial interval match** (*na částečnou intervalovou shodu*)
 - for selected attributes the interval of values entered

File organization

- File organization (*souborová organizace*)
 - how to organize a set of records in a file and how to access them
 - the description of the **logical memory structure** together with **algorithms** for handling that structure
- Files in the organization
 - FO can contain multiple files
- Optimal choice of organization depends on the usage
 - accessing one record, accessing multiple records, accessing whole file, ...

File organization levels

- **Logical schema**

- **algorithms**
 - defined to secure optimal manipulation with the data for given task
 - minimization of the number of operations while manipulating the file
- **logical blocks (pages)**
 - logical blocks structure
 - logical blocks relations
 - logical blocks manipulation
 - fill factor
- **logical files**
 - how are logical pages related to each other
 - **primary file**

- data
- **auxiliary files**
 - efficient access to the data (indexes, ...)

Unlike in case of in-memory data structure, even the structure itself needs to be stored somewhere.

- **Physical schema**

- **mapping** between logical schema and physical pages
- physical files
 - One logical file can span multiple physical files and the other way around

- **Implementation schema**

- **implementation** of the physical files
- **shielded** from the logical level **by OS**

File organization types (1)

- **Heap file** (*halda*)
 - **variable-length** records
 - a record **placed** always at the **end of the file**
- **Sequential file** (*sekvencní soubor*)
 - **unsorted**
 - as heap file but contains **fixed-length** records
 - **sorted**
 - stores records in **sequential order**, based on the **value of the search key** of each record

File organization types

- **Indexed sequential file** (*index-sekvenční soubor*)
 - records stored based on the order of the search key on which index is built
- **Index file** (*indexový soubor*)
 - resembles indexed sequential file but multiple indexes can be present
- **Hashed file** (*hashovaný/ hešovaný soubor*)
 - a hash function is computed on a chosen attribute of a record; the resulting structure specifies in which block of the primary file given record should be placed/found

Heap File (HF)

- **Data not sorted** in any way
 - File is not homogeneous, i.e. contains records of various types/lengths
 - Usually used only along with another supporting structure
- We consider b as an average blocking factor (the records are variable-length)
- **INSERT**
 - fetch the last block in the file and append the new record
 - **$O(1)$**
 - realtime \rightarrow **$O(1) (s+r+btt)$**
 - **FIND**
 - whole file needs to be scanned if the search key attribute set is not unique
 - **$O(N) \rightarrow O(N/b)$**
 - realtime
 - worst case: **$O(N/b)(s+r+btt)$**
 - best case: **$s+r+O(N/b) btt$**

HF – reasons for variable-length records

- When similar objects relate to a non-uniform set of data
- Examples
 - different types of employees within our application → a need to store different information for different roles
 - some records may have non atomic data types, e.g., an unspecified number of repeating groups of attribute values → variable number of items in orders
- Solution
 - variable-length records
 - setting maximum length for a given field

Unsorted sequential file (USF) organization

- Data not sorted in any way
- Suitable when data are collected without any relationship to other data
- INSERT
 - fetch the last block in the file and append the new record
 - $O(1)$
- FETCH
 - Single record
 - whole file needs to be scanned if the search key attribute set is not unique
 - $O(N) \rightarrow O(N/b)$
 - realtime
 - worst case: $O(N/b)(s+r+btt)$
 - best case: $s+r+O(N/b) btt$

Block	Name	Department	...
0	Galvin Janice	Purchasing	
	Walters Rob	Marketing	
	Brown Kevin	Marketing	
1	Walters Rob	Development	
	Duffy Terri	Research	
	Brown Kevin	PR	
2	Duffy Terri	Development	
	Walters David	Production	
	Brown Kevin	Purchasing	
3	Matthew Gigi	Purchasing	
	Walters Rob	PR	
	

Sorted sequential file (SSF) organization

- **Records sorted** in the file on the primary search key
- File can be sorted only according to **one attribute** → the **most often searched** one (if querying is the prevalent operation)

Block	Name	Department	...
0	Brown Kevin	PR	
	Brown Kevin	Purchasing	
1	Brown Kevin	Marketing	
	Duffy Terri	Development	
2	Duffy Terri	Research	
	Galvin Janice	Purchasing	
3	Matthew Gigi	Purchasing	
	Walters David	Production	
4	Walters Rob	Marketing	
	Walters Rob	Development	
5	Walters Rob	PR	
	

SSF – FETCH

- **Sequential scan**

- takes linear time - on average $N/2$ records ($N/b/2$ blocks) need to be checked
- $O(N) \rightarrow O\left(\frac{N}{b}\right)$
 - However, if fetching a range containing k records, only the first record needs to be found, the rest are fetched sequentially $\rightarrow \frac{N/b}{2} + \frac{k}{b} \rightarrow \frac{N/b}{2}(s+r) + \left(\frac{N/b}{2} + \frac{k}{b}\right)btt$

- **Binary search** (*půlení intervalu*)

- $O(\log_2(N)) \rightarrow O(\log_2(N/b)) \rightarrow O(\log_2(N/b))(s+r+btt)$
- address of the i -th block can be obtained from the header information (the file is homogeneous)

SSF - modification

- **INSERT**

- inserting new record into the structure would be costly since all the following records would have to be shifted
- **auxiliary file/blocks** called **overflow file/bucket** (*stránka/ oblast přetečení*) needs to be established where the new records are inserted
- the file is **periodically reorganized**

- **UPDATE**

- simple if the update does not include the primary search key

- **DELETE**

- deleted records are not directly removed since reorganization would have to take place
- **a bit designating deleted records is set**
- deleted records are removed during periodical reorganization

SSF - size

- SSF storing employee records - *Employee (Name, Department, Phone, SSN, ...)*
- Record size $R = 500B$
- Block size $B = 8KiB = 8,192B$
- Records count = 50,000
- Block factor $b = \left\lfloor \frac{B}{R} \right\rfloor = 16$
 - one page can accommodate up to 16 employees
- File size
 - $\frac{50,000}{16} = 3,125$ blocks
 - $3,125 * 8,192 = 25,6 MB$

Indexed sequential file (ISF) organization

- **Primary file/area** (*primární soubor*)
 - data file sorted based on the search key
- **Index/secondary file/area** (*indexový soubor*)
 - primary index
- **Overflow file/area** (*oblast přetečení*)
- Data can be accessed either using the index or sequentially
 - favorable when high percentage of the records is to be fetched

Index

- An **index** is an **auxiliary structure** for a data file that consists of a specifically arranged structure containing **key-pointer pairs**
 - a key-pointer pair might be, e.g., a *Name* value and a disk-address pointer that points to the block that accommodates a record with the given *Name* value
 - a structure containing these pairs would form an index
- **Storage of the index**
 - **main memory**
 - **secondary memory**
 - if the index does not fit in the main memory (can't be cached) it has to be accessed in the same manner as the primary file
 - accessing index must also be taken into account when computing the find/fetch time

ISF – single-level index

- Constant-width records

- Index file associates a key (*Name*) with a pointer to the **block** containing that key

Index file

Block	Name	
6	Brown	0
	Clinard	1
	Duffy	2
	Leavy	3
	Peagler	4
7	Walters Rob	5
	...	

Primary file

Block	Name	Department	...
0	Brown Kevin	PR	
	Berkman Doloris	Purchasing	
1	Clinard Stephnie	Marketing	
	Coolidge Emily	Development	
2	Duffy Terri	Research	
	Galvin Janice	Purchasing	
3	Leavy Shirleen	Purchasing	
	Matthew Gigi	Production	
4	Peagler David	Marketing	
	Shackelford Elsie	Development	
5	Walters Rob	PR	
	

ISF – multiple-level index

- Index is typically **hierarchical** so that **less accesses** are needed
- Master level uses to be stored in the main memory
- In real use, the **blocking factor of the index** tends to be **much higher** than the one of the primary file

Index file 2. (top/master) level		
Block	Name	
8	Brown	6
	Walters Rob	7

Index file 1. (base) level		
Block	Name	
6	Brown	0
	Clinard	1
	Duffy	2
	Leavy	3
	Peagler	4
7	Walters Rob	5
	...	

Primary file		
Block	Name	Department
0	Brown Kevin	PR
	Berkman Doloris	Purchasing
1	Clinard Stephnie	Marketing
	Coolidge Emily	Development
2	Duffy Terri	Research
	Galvin Janice	Purchasing
3	Leavy Shirleen	Purchasing
	Matthew Gigi	Production
4	Peagler David	Marketing
	Shackelford Elsie	Development
5	Walters Rob	PR

ISF – FETCH

- Searching for a **specific value** (query key)
 - check the top level of the index and identify a ***key-value* pair with the highest value lower** than the query key
 - **fetch** the **block** referenced by the *value*
 - **repeat** the previous steps with lower index levels **until a primary file block is reached**
 - **search the primary file block** for the specified key
- Searching for a **range of values**
 - **search** for the **lower bound key of the interval**
 - **sequentially scan** the blocks of the primary file until record corresponding to the upper bound key is found

ISF – size & time (1)

- **Fetch time** depends on the index **tree height**
- **height** = $\lceil \log_p N/b \rceil$, $b = \left\lfloor \frac{B}{R} \right\rfloor$, B = block size, R = (average) record size, $p = \left\lfloor \frac{B}{K+P} \right\rfloor$ block factor for index records (K = key size, P = pointer size)
 - **number of disk accesses** corresponds to the **number of index levels**
 - number of index **levels** can be **decreased** by
 - **decreasing K** (key size)
 - **increasing B** (block size)

ISF – size & time (2)

- ISF storing employee records - **Employee (Name, Department, Phone, SSN, ...)**
 - Let the name of the employee (30B) be the key
- **$R = 500 \text{ B}$, $B = 8 \text{ KiB} = 8,192 \text{ B}$, $b = 16$** , records count = 1,000,000
- **Primary file size**
 - $\lceil 1,000,000/16 \rceil = 62,500$ blocks
 - $62,500 * 8,192 = 512 \text{ MB}$
- **Index height**
 - $p = 215$
 - $\text{height} = \lceil \log_{215} 1000000/16 \rceil = 3$
- **Index size**
 - key size = 30 B, value (pointer) size = 8 B, b for index file = 215
 - $\lceil 62,500/215 \rceil = 290$ blocks → we will need **2 additional blocks in the next level** and **1 block in the top level**
 - $293 * 8,192 \cong 2.4 \text{ MB}$

ISF - update

- **Index structure** stays **static** when inserting data → **new records** need to be stored in **reserved areas** within the primary file
 - when an index is created the index nodes are fixed and do not change during modifications of the primary file
 - **pockets/buckets**
 - each record/block contains a pocket where the overflowed records are stored

ISF – pockets/buckets

- **Overflowed data** are inserted into a **new block** (created **dynamically**) pointed to by the overflowed block
- **Buckets can be chained** and therefore theoretically the ISF does not need to be rebuilt
 - however **long pockets decrease efficiency**
- **Pointers** to the overflow area
 - **with each record**
 - takes more space
 - shorter sequences in the overflow area
 - **with each block**
 - longer sequences in the overflow area
 - it is possible to reserve some space in the block

ISF – pros and cons

Pros

- Fast access using primary search key
- Shares pros of the sequential file

Cons

- Fast access **only** when using primary search key
- Problems with primary file when updating
 - pockets – slows down access to the data
- Possible need of reorganization
 - time consuming operation

Indexed file organization (IF)

- **Motivation**

- being able to search the file according to different attributes without the need to sequentially scan the whole file
 - in the ISF organization this is possible only for one attribute

- **Implementation**


- the **primary file** stays **unsorted** or is **sorted** according to one key only (**primary index**)
- **for each query key** an **index** file can be built → one primary data file, **multiple index files**
 - indexes can be of various types (even within one IF)
- IF basically **corresponds to** a standard DB **table** where we have **one table** and **multiple indexes** built over it (possibly of different types)

Primary vs secondary Index

- **Primary index** (*primární index*)
 - index over **the attribute** based on which the **records in the primary file are sorted**
 - if the value of the primary attribute is modified the file needs to be reorganized → should be relatively invariable
 - well-suited for range queries
 - there does not have to be a primary index in the IF
- **Secondary index** (*sekundární index*)
 - IF can have **multiple** secondary indexes
 - **range queries for long ranges can be very expensive** (an extreme example is a sequential scan based on a secondary index which can lead to an extremely deteriorated performance)

Direct vs. indirect indexing/addressing

- **Direct indexing** (*přímé indexování*)
 - index is **bound directly to the record**
 - primary file reorganization leads to modification of all the indexing structures
- **Indirect indexing** (*nepřímé indexování*)
 - **secondary indexes** contain **keys of the primary index** and not pointers to the primary file
 - accessing a record needs one more accesses to the primary index
 - if the primary file is reorganized, the secondary indexes stay intact



IF – example (1)

- Unsorted primary file
- Records in the index files are sorted based on given key

			Primary file			
Blo ck	Department		Blo ck	Name	Department	...
33	Development	24	24	Development	1.1	0 Brown Kevin PR
	Marketing	25		Development	4.1	Berkman Doloris Purchasing
	Production	26		Marketing	1.0	1 Clinard Stephnie Marketing
34	Purchasing	27	25	Marketing	4.0	Coolidge Emily Development
				PR	0.0	2 Duffy Terri Research
				PR	5.0	Galvin Janice Purchasing
			26	Production	3.1	3 Leavy Shirleen Purchasing
				Purchasing	0.1	Matthew Gigi Production
				Purchasing	2.1	4 Peagler David Marketing
			27	Purchasing	3.0	Shackelford Elsie Development
				Research	2.0	5 Walters Rob PR
		

Block ID

Offset

IF – example (2)

- Primary file sorted based on the primary key “Name”

Blo ck	Name	
28	Brown	24
	Matthew	25
	Young	26
29	Zhang	27

Blo ck	Name	
24	Brown	0
	Clinard	1
	Duffy	2
25	Leavy	3
	Peagler	4
	Walters Rob	5
26	Young	6
	...	
27	Zhang	

Primary file			
Blo ck	Name	Department	...
0	Brown Kevin	PR	
	Berkman Doloris	Purchasing	
1	Clinard Stephnie	Marketing	
	Coolidge Emily	Development	
2	Duffy Terri	Research	
	Galvin Janice	Purchasing	
3	Leavy Shirleen	Purchasing	
	Matthew Gigi	Production	
4	Peagler David	Marketing	
	Shackelford Elsie	Development	
5	Walters Rob	PR	
...	

IF – example (3)

- If the primary file is sorted on the the index key, addressing blocks in the primary file is sufficient
 - Less entries in the leaf level of the index than the actual records
- If the primary file is not sorted on the index key, leaf level needs to address every record in the primary file
 - E.g. department records are spread all over the primary file

IF – example (4)

- Primary index “Name” and secondary index “Department”, primary file sorted over primary index

Secondary index		
Blo ck	Department	
8	Development	Coolidge
	Development	Shackelford
	Marketing	Clinard
9	Marketing	Peagler
	PR	Brown
	PR	Walters
10	Production	Matthew
	Purchasing	Galvin
	Purchasing	Leavy
11	Research	Duffy

Primary index		
Blo ck	Name	
6	Brown	0
	Clinard	1
	Duffy	2
7	Leavy	3
	Peagler	4
	Walters Rob	5
	...	

Primary file			
Blo ck	Name	Department	...
0	Brown Kevin	PR	
	Berkman Doloris	Purchasing	
1	Clinard Stephnie	Marketing	
	Coolidge Emily	Development	
2	Duffy Terri	Research	
	Galvin Janice	Purchasing	
3	Leavy Shirleen	Purchasing	
	Matthew Gigi	Production	
4	Peagler David	Marketing	
	Shackelford Elsie	Development	
5	Walters Rob	PR	

IF – comments

- The **indirect indexing** is in nowadays databases often the method of choice
- It is desirable for the **primary index** to stay **in main memory** and not to be swapped to disk
- **Keys** of a **primary index** are advised to be **small**, e.g., integers and not long strings

Hashed file organization

- **Direct access**
 - suitable when **accessing individual records** given **one unique key**
 - uses a **hash function** to **map keys to** the block/page **addresses**
 - **hashing field**
 - the attribute over which the quick access is provided
 - typically **$O(1)$** time
 - if the data can not fit into the page when inserting, an overflow strategy is employed

HF - example

- Hash function h uses three lower bits of the first integer field to address a page where the record is stored
 - $h(\{42, \text{true}, \text{"foo"}\}) \rightarrow 2$ ($42 = 101010_2$)
 - $h(\{14, \text{true}, \text{"bar"}\}) \rightarrow 6$ ($14 = 1110_2$)
 - $h(\{26, \text{false}, \text{"false"}\}) \rightarrow 2$ ($26 = 11010_2$)
- Placement within the page is not specified
- When the file is being reorganized, the pages are filled to only ,e.g., 80%
 - avoiding an overflow immediately after the structure is built

Operating instructions

