

Data Organization and Processing

Spatial Join

(NDBI007)

David Hoksza, Škoda Petr
<http://siret.ms.mff.cuni.cz/hoksza>

Outline

- Spatial join basics
- Relational join
- Spatial join

Spatial join definition (1)

- Given **two sets of multi-dimensional** objects in Euclidean space, a spatial join finds **all pairs** of objects **satisfying** a given **spatial relation** between the objects, such as intersection.
- **Simplified spatial join**
 - Given two sets of rectangles, R and S , find all of the pairs of intersecting rectangles between the two sets, i.e. $\{(r, s) : r \cap s \neq \emptyset, r \in R, s \in S\}$

Spatial join definition (2)

- **Spatial overlay join (general spatial join)**
 - the data set can consist of **general spatial objects** (points, lines, polygons)
 - the data sets can have **more** than **two dimensions**
 - the relation between pairs of spatial objects can **be any spatial relation** (nearness, enclosure, direction, ...)
 - there can be more sets in the relation (**multiway spatial join**) or one set joined with itself (**self spatial join**)

Spatial join examples

- Find all pair of rivers and cities that intersect.
- Find all of the rural areas that are below sea level, given an elevation and land use map.
- Find the houses inside the areas with poor slope stability.

Non-spatial join

- There exist algorithms for **standard relational join**
 - only equi-joins (the join predicate is equality) considered here
 - **nested loop** join (*hnížděné cykly*)
 - **sort-merge** join (*setřídění-slévání*)
 - **hash** join (*hashované spojení*)
- most of the standard relational join algorithms are not suitable for spatial data because the join condition involves multidimensional spatial attribute

Non-spatial nested loop join


- Nested loop join checks one by one for each element of a relation R all elements in relation in S

```
FOREACH  $r \in R$  DO
  FOREACH  $s \in S$  DO
    IF  $\text{cond}(r, s)$  THEN REPORT( $r, s$ )
```

- In its basic version, the nested loop join is the least efficient algorithm from the relational joins algorithms
- The only of the standard relational join algorithms applicable also to spatial data
 - Therefore special spatial join algorithms had to be developed

Non-spatial sort-merge join

- Two-phase algorithm
 - sort both relations R, S independently
 - scan both relations at once in the same order and join
- Multi-dimensional data do not preserve proximity so this method (used as is) is not applicable to spatial data



Non-spatial hash join



- Hash join
 - One of the relations (R) is hashed with a hash function h (the assumption that R fits into main memory)
 - The other relation (S) is processed one by one and the elements' ids are hashed with h
 - If for two elements $r \in R, s \in S: h(r) = h(s)$ then r and s are checked for $r.cmpr_attributes = s.cmpr_attributes$
- Equijoins rely on grouping objects with the same value which is not possible for spatial objects since these have an extent

Filter-refine policy


- The **objects** to be searched for **can be very complex**
 - → testing of the join condition (spatial predicate) itself can be highly time demanding
 - → not many objects fit into main memory
 - → filter-refine strategy
- Spatial objects are approximated using simple spatial objects
- **Filter-refine**
 - **filter** - the spatial join is conducted using the objects' **approximations** → **candidate set**
 - **refine** – pairs which pass the filter are **tested for the spatial predicate** using their **full spatial representation**

Approximations (1)

- The most common approximation is **minimum bounding rectangle (MBR)** – the smallest rectangle fully enclosing given object whose sides are parallel to the axes
- **Dead space**
 - the amount of **space covered by the approximation but not the approximated object**
 - the approximation should aim to **minimize dead space**
 - large dead space areas can lead to higher false hit rate in the filtering stage



Large dead area



False hit

Approximations (2)


- **Conservative**

- every point of the object is also point of the approximation
 - concave
 - convex
 - MBR – most common
 - minimum bounding polygon
 - minimum bounding circle/ellipse



- **Progressive**


- every point of the approximation is also point of the object
 - maximum nested rectangles, circles, ...



Approximations (3)

- **Object decomposition**

- minimizes dead space by decomposing an object into disjoint fragments with their own MBRs
- after refine step and extra filtering stage is needed to remove duplicities



Spatial join methods

Internal Memory Methods

- Nested loop join
- Index nested loop join
- Plane Sweep
- Z-order

External Memory Methods

- Hierarchical traversal
- Partitioning-based methods

Internal memory methods

- Nested loop join
- Index nested loop join
- Plane Sweep
- Z-order

Nested loop join

- The algorithm is **identical to the standard relational one**
 - works with **arbitrary object type** and **join condition**
- Given datasets A and B the join takes $O(|A| \times |B|)$ time
 - suitable for **small datasets**

NESTED_LOOP_JOIN(setA, setB, joinCondition)

INPUT: Sets to join and condition based on which the join happens

OUTPUT: Pairs of objects satisfying the join condition

FOREACH $a \in \text{setA}$

 FOREACH $b \in \text{setB}$

 IF Satisfied($a, b, \text{joinCondition}$) THEN

 REPORT(a, b);

Index nested loop join

- Variant of nested loop join where **first a spatial index** is created over one of the sets
- The **indexed set is queried** by each of the objects from the second set for intersection (window query)
- Given datasets A and B the join takes $O(\log(|A|) \times |B|)$ time (not including time needed for building the index)

INDEX_NESTED_LOOP_JOIN(setA, setB)

INPUT: Sets to join

OUTPUT: Pairs of intersecting objects


```
ix := CreateSpatialIndex(setA)
```

```
FOREACH b ∈ setB
```

```
    REPORT (ix.Search(b)) ;
```

Plane sweep

- **Two-phase** algorithm for identification of intersecting rectangles
 - **sorting the rectangles** in ascending order on the basis of their left sides (x -axis)
 - **sweeping a vertical scan line** through the sorted list from left to right, halting at each of the rectangles lower x coordinates and identification of rectangles intersecting (vertical line) with the current rectangle and checking for intersection based on the y -axis



Similar to non-spatial
sort-merge join


Plane sweep algorithm (1)

PLANE_SWEEP(setA, setB)

INPUT: Sets of rectangles to join

OUTPUT: Pairs of intersecting rectangles

```
listA ← SortByLeftSide(setA);  
listB ← SortByLeftSide(setB);  
sweepStructA ← CreateSweepStructure();  
sweepStructB ← CreateSweepStructure();  
WHILE NOT(listA.End()) || NOT (listB.End()) DO  
    IF listA.First() < listB.First() THEN  
        sweepStructA.Insert(listA.First());  
        sweepStructB.RemoveInactive(listA.First());  
        sweepStructB.Search(listA.First());  
        listA.Next();  
    ELSE  
        sweepStructB.Insert(listB.First());  
        sweepStructA.RemoveInactive(listB.First());  
        sweepStructA.Search(listB.First());  
        listB.Next();
```



Plane sweep algorithm (2)

- Sweep structure tracks active rectangles and has to support three operations
 - **Insert**
 - insert a rectangle into the active set
 - **RemoveInactive**
 - removes from the active set all rectangles that do not overlap a given rectangle (line)
 - **Search**
 - searches for all active rectangles that intersect a given rectangle and outputs them
 - if the data are sorted, only the data in the sweep structures need to be kept in internal memory
- Given datasets A and B the algorithm takes $O(\log(n) \times n)$ where $n = |A + B|$, including the list sort time.

Z-ordering

- Z-order methods are based on **representing an object using a set of continuous segments corresponding to Z-ordering**
 - object $o \rightarrow \{(z_1, o), \dots (z_n, o)\}$
- The z-order-based representation can be processed in different ways
 - **sort-merge**
 - standard relational sort-merge but the sets are z-ordered
 - **plane-sweep**
 - the active set, instead of being the objects that intersect the sweep line, are the enclosing cells of the objects that intersect the current Z-order grid cell
 - the sweep structure can be represented by a stack where each object in the stack intersects the query object

External memory methods

- Hierarchical traversal
- Partitioning-based methods

Hierarchical traversal

- Assumption - **both datasets** are **indexed** using a hierarchical index such as R-tree
- **Synchronized traversal** can be used to test the join condition
 - the algorithm traverses the two trees in a synchronized fashion and **compares bounding objects at given levels**
 - for indexes of different heights, the join of a leaf of one index with a sub-tree of the other can be accomplished using a window query, or handling leaf-to-node comparison as special cases
 - if a node corresponding to a part of the space does not match the condition it can be excluded from the traversal
 - iterative filter and refine approach
 - reading a node usually corresponds to reading one memory page

General synchronized traversal

INDEXED_TRAVERSAL_JOIN(rootA, rootB)


INPUT: Roots of the structures representing the sets to be joined

OUTPUT: Pairs of intersecting rectangles

```
queue ← CreateQueue();
queue.Add(pair(rootA, rootB));
WHILE NOT(queue.Empty()) DO
    nodePair ← queue.Pop();
    pairs ← IdentifyIntersectingPairs(nodePair);
    FOREACH p ∈ pairs DO
        IF p is leaf THEN ReportIntersection(p);
        ELSE queue.Add(p);
```


Partitioning

- Often applied when neither of the sets to be joined is indexed
- Resulting partitions should be small enough to fit in internal memory
- Once the data are partitioned, each pair of overlapping partitions is read into internal memory and internal memory techniques are used



Partition join of uniform data

GRID_JOIN(setA, setB)

INPUT: Sets of objects to be joined

OUTPUT: Pairs of intersecting objects


```
m ← AvailableInternalMemory();
mbrSize ← BytesToStoreMBR();
minNrOfPartitions ← (setA.Size() + setB.size()) * mbrSize() / m;
partList ← DeterminePartitions(minNrOfPartitions);

{object appears in every partition it intersects}
partitionPointersA ← PartitionData(partList, SetA);
partitionPointersB ← PartitionData(partList, SetB);

FOREACH part ∈ partList DO
    partitionA ← ReadPartition(partitionPointersA, part);
    partitionB ← ReadPartition(partitionPointersB, part);
    PLANE_SWEEP(partitionA, partitionB);
```

Avoiding duplicate results

- Space partitioning can lead to object duplication which can, in turn, lead to duplication of results



- **Deduplication**

- **sort and remove duplicates**



- requires sorting which implies increased computational demands

- **reference point method**

- a consistently chosen reference point is selected from the intersecting region
 - intersection is reported only if the reference point lies within given partition


Grid partitioning – skewed distributions

- Basic grid algorithm is rarely used since the objects distribution is often not uniform
- [Patel & DeWitt; 1996] proposed to group partitions using a mapping function to minimize skew by creating partitions having similar number of items



Size-based partitioning

- [Koudas & Sevcik; 1997]
- Partitioning of data using a **series of finer and finer grids**
 - **level i grid contains 4^{i-1} partitions** (i horizontal and i vertical splitting axis)
 - **each object is placed into the partition** associated with the grid in which the object **does not intersect the grid lines**
 - in second-level grid are objects which cross level 3 partition borders (16 cells) but not level 2 partition borders (4 cells)
 - each partition is a filter and objects fall through to the lowest level partition where a partition boundary is crossed



Searching in size-based partitioning

- Given two spatial data sets A and B
 - Scan data sets A and B and for each entity:
 1. Compute the Hilbert (Z) value of the entity
 2. Determine the level at which the entity belongs and place its entity descriptor in the corresponding level file
 - For each level file
 1. Sort by Hilbert (Z) value
 - Perform a synchronized scan over the pages of level files
- Each page of each level file is read just once
 - Entries in page A^l (l -th level) need to be checked only against actual pages B^m ($m \leq l$)

